

In [1]:

```
#Importing the libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

import sklearn
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

from scipy.cluster.hierarchy import linkage
from scipy.cluster.hierarchy import dendrogram
from scipy.cluster.hierarchy import cut_tree
```

In [2]:

```
#Loading the data into sales as pandas dataframe
sales = pd.read_excel("sales_data.xlsx")
sales.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 5000 entries, 0 to 4999

Data columns (total 40 columns):

#	Column	Non-Null Count	Dtype
0	CustomerID	5000 non-null	int64
1	TOTAL_ORDERS	5000 non-null	int64
2	REVENUE	5000 non-null	float64
3	AVERAGE_ORDER_VALUE	5000 non-null	float64
4	CARRIAGE_REVENUE	5000 non-null	float64
5	AVERAGESHIPPING	5000 non-null	float64
6	FIRST_ORDER_DATE	5000 non-null	datetime64[ns]
7	LATEST_ORDER_DATE	5000 non-null	datetime64[ns]
8	AVGDAYSBEETWEENORDERS	5000 non-null	float64
9	DAYSSINCELASTORDER	5000 non-null	int64
10	MONDAY_ORDERS	5000 non-null	int64
11	TUESDAY_ORDERS	5000 non-null	int64
12	WEDNESDAY_ORDERS	5000 non-null	int64
13	THURSDAY_ORDERS	5000 non-null	int64
14	FRIDAY_ORDERS	5000 non-null	int64
15	SATURDAY_ORDERS	5000 non-null	int64
16	SUNDAY_ORDERS	5000 non-null	int64
17	MONDAY_REVENUE	5000 non-null	float64
18	TUESDAY_REVENUE	5000 non-null	float64
19	WEDNESDAY_REVENUE	5000 non-null	float64
20	THURSDAY_REVENUE	5000 non-null	float64
21	FRIDAY_REVENUE	5000 non-null	float64
22	SATURDAY_REVENUE	5000 non-null	float64
23	SUNDAY_REVENUE	5000 non-null	float64
24	WEEK1_DAY01_DAY07_ORDERS	5000 non-null	int64
25	WEEK2_DAY08_DAY15_ORDERS	5000 non-null	int64
26	WEEK3_DAY16_DAY23_ORDERS	5000 non-null	int64
27	WEEK4_DAY24_DAY31_ORDERS	5000 non-null	int64
28	WEEK1_DAY01_DAY07_REVENUE	5000 non-null	float64
29	WEEK2_DAY08_DAY15_REVENUE	5000 non-null	float64
30	WEEK3_DAY16_DAY23_REVENUE	5000 non-null	float64
31	WEEK4_DAY24_DAY31_REVENUE	5000 non-null	float64
32	TIME_0000_0600_ORDERS	5000 non-null	int64
33	TIME_0601_1200_ORDERS	5000 non-null	int64
34	TIME_1200_1800_ORDERS	5000 non-null	int64
35	TIME_1801_2359_ORDERS	5000 non-null	int64
36	TIME_0000_0600_REVENUE	5000 non-null	float64

```

37 TIME_0601_1200_REVENUE      5000 non-null    float64
38 TIME_1200_1800_REVENUE      5000 non-null    float64
39 TIME_1801_2359_REVENUE      5000 non-null    float64
dtypes: datetime64[ns](2), float64(20), int64(18)
memory usage: 1.5 MB

```

```

In [3]: #sorting the data according to 'CustomerID'
sales = sales.sort_values(by=['CustomerID'])
sales.head()


```

```

Out[3]:
   CustomerID  TOTAL_ORDERS  REVENUE  AVERAGE_ORDER_VALUE  CARRIAGE_REVENUE  AVERAGESHIPPING
2266         1             61  34847.40                571.27              297.50             4.88
2876         2             59  32486.98                550.63              218.68             3.71
2267         3             53  24178.97                456.21               43.97             0.83
1153         4             84  18554.49                220.89              421.29             5.02
3377         5             26  16884.99                649.42               54.89             2.11

```

5 rows × 40 columns



RFM ANALYSIS Since we're performing an RFM analysis thus we need to consider only the Recency (How recently the orders were made), Frequency(How frequently a particular customer made the order or the total number of orders made by the customer in the given time period) and Monetary value (What was the avg amount of order made by customer).

```

In [4]: #FREQUENCY
frequency = sales[['CustomerID', 'TOTAL_ORDERS']]
frequency.head()

```

```

Out[4]:
   CustomerID  TOTAL_ORDERS
2266         1             61
2876         2             59
2267         3             53
1153         4             84
3377         5             26

```

```

In [5]: frequency = frequency.reset_index()
frequency.drop(['index'], axis=1, inplace=True)
frequency.head()

```

```

Out[5]:
   CustomerID  TOTAL_ORDERS
0           1             61
1           2             59
2           3             53
3           4             84
4           5             26

```

```

In [6]: #MONETARY
groupDf = sales[['CustomerID', 'AVERAGE_ORDER_VALUE']]

```

```
groupDf = groupDf.reset_index()
groupDf.drop(['index'],axis=1,inplace=True)
groupDf.head()
```

Out[6]:

	CustomerID	AVERAGE_ORDER_VALUE
--	------------	---------------------

0	1	571.27
1	2	550.63
2	3	456.21
3	4	220.89
4	5	649.42

In [7]:

```
#Merging Frequency and Monetary DFs(groupDf)
groupDf = pd.merge(groupDf,frequency, on='CustomerID', how='inner')
groupDf.head()
```

Out[7]:

	CustomerID	AVERAGE_ORDER_VALUE	TOTAL_ORDERS
--	------------	---------------------	--------------

0	1	571.27	61
1	2	550.63	59
2	3	456.21	53
3	4	220.89	84
4	5	649.42	26

In [8]:

```
groupDf.columns = ['CustomerID','amount','frequency']
groupDf.head()
```

Out[8]:

	CustomerID	amount	frequency
--	------------	--------	-----------

0	1	571.27	61
1	2	550.63	59
2	3	456.21	53
3	4	220.89	84
4	5	649.42	26

In [9]:

```
#RECENCY
recency = sales[['CustomerID','DAYSSINCELASTORDER']]
recency = recency.reset_index()
recency.drop(['index'],axis=1,inplace=True)
recency.head()
```

Out[9]:

	CustomerID	DAYSSINCELASTORDER
--	------------	--------------------

0	1	53
1	2	94
2	3	53
3	4	5
4	5	130

```
In [10]: #Merging groupDf and RECENCY dfs to get a consolidated one
groupDf = pd.merge(groupDf, recency, on='CustomerID', how='inner')
groupDf.head()
```

```
Out[10]:
```


	CustomerID	amount	frequency	DAYSSINCELASTORDER
0	1	571.27	61	53
1	2	550.63	59	94
2	3	456.21	53	53
3	4	220.89	84	5
4	5	649.42	26	130

```
In [11]: groupDf.columns = ['CustomerID', 'amount', 'frequency', 'recency']
groupDf.head()
```


```
Out[11]:
```

	CustomerID	amount	frequency	recency
0	1	571.27	61	53
1	2	550.63	59	94
2	3	456.21	53	53
3	4	220.89	84	5
4	5	649.42	26	130

```
In [12]: #checking the data for further clustering prcessing and analysis
plt.boxplot(groupDf['recency'])
plt.show()
```




```
In [13]: plt.boxplot(groupDf['frequency'])
plt.show()
```



In [14]:

```
plt.boxplot(groupDf['amount'])
plt.show()
```



Since there are a lot of outliers in the data, in frequency and amount. Thus proceeding with outlier's treatment.

Outlier's Treatment

In [15]:

```
# for amount
q1 = groupDf.amount.quantile(0.05)
q3 = groupDf.amount.quantile(0.95)
iqr = q3-q1
groupDf = groupDf[(groupDf.amount >= q1 - 1.5*iqr)&(groupDf.amount <= q3+1.5*iqr)]

# for frequency
q1 = groupDf.frequency.quantile(0.05)
q3 = groupDf.frequency.quantile(0.95)
iqr = q3-q1
groupDf = groupDf[(groupDf.frequency >= q1 - 1.5*iqr)&(groupDf.frequency <= q3+1.5*iqr)]

# for recency
q1 = groupDf.recency.quantile(0.05)
q3 = groupDf.recency.quantile(0.95)
iqr = q3-q1
groupDf = groupDf[(groupDf.recency >= q1 - 1.5*iqr)&(groupDf.recency <= q3+1.5*iqr)]
groupDf.head()
```

Out[15]:

	CustomerID	amount	frequency	recency
0	1	571.27	61	53

	CustomerID	amount	frequency	recency
1	2	550.63	59	94
2	3	456.21	53	53
3	4	220.89	84	5
4	5	649.42	26	130

```
In [16]: # Rescaling the data
rfmDf = groupDf[['amount', 'frequency', 'recency']]

# instantiating
scaler = StandardScaler()

# Using Fit-Transform

rfmDfScaled = scaler.fit_transform(rfmDf)
rfmDfScaled.shape
```

Out[16]: (4973, 3)

```
In [17]: rfmDfScaled = pd.DataFrame(rfmDfScaled)
rfmDfScaled.columns = ['amount', 'frequency', 'recency']
rfmDfScaled.head()
```

```
Out[17]:
```

	amount	frequency	recency
0	5.507327	4.025029	-0.430091
1	5.247431	3.858249	0.080965
2	4.058503	3.357910	-0.430091
3	1.095377	5.942994	-1.028399
4	6.491385	1.106386	0.529696

Modelling

Using K-Means Algorithm

```
In [18]: # Using an arbitrary number (n_cluster = 4)
kmeans = KMeans(n_clusters=4, max_iter=50)
kmeans.fit(rfmDfScaled)
```

Out[18]: KMeans(max_iter=50, n_clusters=4)

```
In [19]: kmeans.labels_
```

Out[19]: array([3, 3, 3, ..., 1, 1, 0])

```
In [20]: # sum of squared distance (ssd)
ssd=[]
range_n_clusters= [2,3,4,5,6,7,8,9]
for num_clusters in range_n_clusters:
    kmeans = KMeans(n_clusters = num_clusters, max_iter=50)
    kmeans.fit(rfmDfScaled)
```


```

ssd.append(kmeans.inertia_)

plt.plot(ssd)

```

Out[20]: [



Hopkins Test

Going for Hopkins Test to check if the data is clusterable

In [21]:

```

from sklearn.neighbors import NearestNeighbors
from random import sample
from numpy.random import uniform
import numpy as np
from math import isnan

def hopkins(X):
    d = X.shape[1]
    #d = len(vars) # columns
    n = len(X) # rows
    m = int(0.1 * n)
    nbrs = NearestNeighbors(n_neighbors=1).fit(X.values)

    rand_X = sample(range(0, n, 1), m)

    ujd = []
    wjd = []
    for j in range(0, m):
        u_dist, _ = nbrs.kneighbors(uniform(np.amin(X,axis=0),np.amax(X,axis=0),d).reshape(1, -1))
        ujd.append(u_dist[0][1])
        w_dist, _ = nbrs.kneighbors(X.iloc[rand_X[j]].values.reshape(1, -1), 2, return_distance=True)
        wjd.append(w_dist[0][1])

    H = sum(ujd) / (sum(ujd) + sum(wjd))
    if isnan(H):
        print(ujd, wjd)
        H = 0

    return H

```

In [22]:

```
hopkins(rfmDfScaled)
```

Out[22]: 0.9209553084102718

Since the number is close to one thus the data is clusterable. Now proceeding with Silhouette's Analysis.

Silhouette's Analysis

From the elbow curve we found that the X-point (2), i.e., 4 clusters are optimal for the data. Thus using 4 as the optimal number.

```
In [23]: kmeans = KMeans(n_clusters = 4,max_iter=50)
kmeans.fit(rfmDfScaled)
```

```
Out[23]: KMeans(max_iter=50, n_clusters=4)
```

```
In [24]: kmeans.labels_
```


```
Out[24]: array([2, 2, 2, ..., 3, 3, 1])
```

```
In [25]: # Assigning the cluster label to the groupDf
groupDf['clusterID'] = kmeans.labels_
groupDf.head()
```


```
Out[25]:
```

	CustomerID	amount	frequency	recency	clusterID
0	1	571.27	61	53	2
1	2	550.63	59	94	2
2	3	456.21	53	53	2
3	4	220.89	84	5	0
4	5	649.42	26	130	2


```
In [26]: #plotting
sns.boxplot(x='clusterID',y='amount',data=groupDf)
plt.show()
```



```
In [27]: sns.boxplot(x='clusterID',y='frequency',data=groupDf)
plt.show()
```

```
In [28]: sns.boxplot(x='clusterID',y='recency',data=groupDf)
plt.show()
```



```
In [29]: rfmDfScaled.head()
```

```
Out[29]:
```

	amount	frequency	recency
0	5.507327	4.025029	-0.430091
1	5.247431	3.858249	0.080965
2	4.058503	3.357910	-0.430091
3	1.095377	5.942994	-1.028399
4	6.491385	1.106386	0.529696

Hierarchical Clustering

```
In [30]: # viewing the data
rfmDfScaled.head()
```

```
Out[30]:
```

	amount	frequency	recency
0	5.507327	4.025029	-0.430091
1	5.247431	3.858249	0.080965
2	4.058503	3.357910	-0.430091

	amount	frequency	recency
3	1.095377	5.942994	-1.028399
4	6.491385	1.106386	0.529696


In [31]: `groupDf.head()`

Out[31]:

	CustomerID	amount	frequency	recency	clusterID
0	1	571.27	61	53	2
1	2	550.63	59	94	2
2	3	456.21	53	53	2
3	4	220.89	84	5	0
4	5	649.42	26	130	2


In [32]:

```
# performing the Single Linkage
mergings = linkage(rfmDfScaled , method = 'single' , metric = 'euclidean')
dendrogram(mergings)
plt.show()
```



In [33]:

```
# Complete Linkage
mergings = linkage(rfmDfScaled, method = 'complete', metric = 'euclidean')
dendrogram(mergings)
plt.show()
```



```
In [34]: # 4 clusters
cluster_labels = cut_tree(mergings,n_clusters = 4).reshape(-1,)
cluster_labels
```


```
Out[34]: array([0, 0, 0, ..., 3, 3, 3])
```

```
In [35]: groupDf['cluster_labels'] = cluster_labels
groupDf.head()
```


```
Out[35]:
```

	CustomerID	amount	frequency	recency	clusterID	cluster_labels
0	1	571.27	61	53	2	0
1	2	550.63	59	94	2	0
2	3	456.21	53	53	2	0
3	4	220.89	84	5	0	1
4	5	649.42	26	130	2	0

```
In [36]: sns.boxplot(x = 'cluster_labels', y = 'recency', data = groupDf)
plt.show()
```




```
In [37]: sns.boxplot(x = 'cluster_labels', y = 'frequency', data = groupDf)
plt.show()
```



```
In [38]: sns.boxplot(x = 'cluster_labels', y = 'amount', data = groupDf)
```

```
plt.show()
```



As we can see from the above three boxplots, the Cluster 4 (cluster_label = 3) has made the most recent order and it also makes the orders more frequently as compared to others. However when we compare the average amount per order it turns out to be the least for the customers in the cluster 4. We can also consider cluster 3 (cluster_label = 2) for the most recent order after cluster 4 and cluster 3 has made the most frequent orders after cluster 4 however when we compare the average amount per order cluster 3 (cluster_labels = 2) scores a second position again. So we should be focusing on Cluster 3 (cluster_labels = 2) if we want to have a good amount of sales and profits.


We could also see that cluster 1 (cluster_labels = 0) has made the purchase for highest amounts but have made the order less frequently. We could ask our marketing team to roll out more attention grabbing advertisements for those particular set of customers in order to have more revenue as they have also made orders recently which means Cluster 1 customers are loyal too.

We could also proceed with the scatter plot and see which users have made the highest number of orders on each day of the week with the highest revenue. Similarly we could go with the weekwise data and a particular day's time wise data so that we could target the potential customers when there's a high likelihood for them to buy.

For example: For checking customerwise report of orders and revenue on monday we could use a scatter plot.

In [46]:

```
plt.scatter(x=sales['MONDAY_ORDERS'],y=sales['MONDAY_REVENUE'])
plt.title('Monday Order v/s Revenue', color= 'Red')
plt.xlabel('Orders made')
plt.ylabel('Revenue made')
plt.show()
```



In []: